

Final project of operating system

Description

FUSE (<http://fuse.sourceforge.net/>) is a Linux kernel extension that allows for a user space program to provide the implementations for the various file-related syscalls. We will be using FUSE to create our own file system, managed a file that represent our disk device. Through FUSE and our implementation, it will be possible to interact with our newly created file system using standard UNIX/Linux programs in a transparent way.

From a user interface perspective, our file system will be a two level directory system, with the following restrictions/simplifications:

1. The root directory “\” will not only contain other subdirectories, but also regular files
2. The subdirectories will only contain regular files, and no subdirectories of their own
3. All files will be full access (i.e., `chmod 0666`), with permissions to be mainly ignored
4. Many file attributes such as creation and modification times will not be accurately stored
5. Files cannot be truncated
6. Directories are treated as files

From an implementation perspective, the file system will keep data on “disk” via a linked allocation strategy, outlined below.

The kernel version should be 2.6.18, FUSE version should be 2.7.0. You can download them at our website.

Installation of FUSE

The first thing you should do is install source code of 2.6.18 kernel, and rebuild it.

Then install FUSE.

FUSE consists of two major components: A kernel module that has already been installed, and a set of libraries and example programs that you need to install.

First, copy the source code to your directory

```
tar xvf fuse-2.7.0.tar
```

```
cd fuse-2.7.0
```

Now, we do the normal configure, compile, install procedure on UNIX, but omit the install step since that needs to be done as a superuser and has already been done.

```
./configure
```

```
make
```

(The third step would be `make install`, but if you try it, you will be met with many access denied errors.)

First FUSE Example

Let us now walk through one of the examples. Enter the following:

```
cd fuse-2.7.0/examples
```

```
mkdir testmount
```

```
ls -al testmount
```

```
./hello testmount
```

```
ls -al testmount
```

You should see 3 entries: `.`, `..`, and `hello`. We just created this directory, and thus it was empty, so where did `hello` come from? Obviously the `hello` application we just ran could have created it, but what it actually did was lie to the operating system when the OS asked for the contents of that directory. So let's see what happens when we try to display the contents of the file.

```
cat testmount/hello
```

You should get the familiar hello world quotation. If we `cat` a file that doesn't really exist, how do we get meaningful output? The answer comes from the fact that the `hello` application also gets notified of the attempt to read and open the fictional file "`hello`" and thus can return the data as if it was really there.

Examine the contents of `hello.c` in your favorite text editor, and look at the implementations of `readdir` and `read` to see that it is just returning hard coded data back to the system.

The final thing we always need to do is to unmount the file system we just used when we are done or need to make changes to the program. Do so by:

```
fusermount -u testmount
```

FUSE High-level Description

The `hello` application we ran in the above example is a particular FUSE file system provided as a sample to demonstrate a few of the main ideas behind FUSE. The first thing we did was to create an empty directory to serve as a mount point. A mount point is a location in the UNIX hierarchical file system where a new device or file system is located. As an analogy, in Windows, “My Computer” is the mount point for your hard disks and CD-ROMs, and if you insert a USB drive or MP3 player, it will show up there as well. In UNIX, we can have mount points at any location in the tree.

Running the `hello` application and passing it the location of where we want the new file system mounted initiates FUSE and tells the kernel that any file operations that occur under our now mounted directory will be handled via FUSE and the `hello` application. When we are done using this file system, we simply tell the OS that it no longer is mounted by issuing the above `fusermount -u` command. At that point the OS goes back to managing that directory by itself.

What You Need to Do

Your job is to create the `u_fs` file system as a FUSE application that provides the interface described in the first section.

The `u_fs` file system should be implemented using an image file, managed by the real file system in the directory that contains the `u_fs` application. The layout of file system will be follow. We will consider the disk to have 512 byte blocks.

Super block (1 block)	Bitmap block (1280 blocks)	Data block (all the rest blocks)
--------------------------	-------------------------------	-------------------------------------

Super block

Super block must be the first block of the file system. It describes the whole file system. Info containing in super block should be:

```
struct sb {  
  
    long fs_size;    //size of file system, in blocks  
  
    long first_blk;  //first block of root directory  
  
    long bitmap;     //size of bitmap, in blocks  
  
}
```

Directories

Directories should be also treated as a file. Each directory contains a list of `u_fs_directory_entry` structures. There is no limit on how many directories we can have.

```
struct u_fs_file_directory    {
    char fname[MAX_FILENAME + 1];    //filename (plus space for nul)
    char fext[MAX_EXTENSION + 1];    //extension (plus space for
    nul)
    size_t fsize;                    //file size
    long nStartBlock;                //where the first block is on
    disk
    int flag;                        //indicate type of file. 0:for unused; 1:for file; 2:for
    directory
}
```

Each directory entry will contain an 8-character maximum directory name, and then have a list of files that are in the directory.

Each file entry in the directory has a filename in 8.3 format. We also need to record the total size of the file, and the location of the file's first block on disk.

Files

Files will be stored in a virtual disk that is implemented as a single, pre-sized file called `.disk` with 512 byte blocks of the format:

```
struct u_fs_disk_block {
    size_t size; // how many bytes are being used in this block
    long nNextBlock; //The next disk block, if needed. This is the next
    pointer in the linked allocation list
    char data[MAX_DATA_IN_BLOCK]; // And all the rest of the space in the block
    can be used for actual data storage.
};
```

Disk Management

In order to manage the free or empty space, you will need to create some bitmap blocks on the disk that record whether a given block has been previously allocated or not. The total number of bitmap blocks should depends on the size of file system. You can do this however you like.

To create a 5MB disk image, execute the following:

```
dd bs=1K count=5K if=/dev/zero of=diskimg
```

This will create a file initialized to contain all zeros, named diskimg. You only need to do this once, or every time you want to completely destroy the disk.

You should also write a format program to init this file, i.e. write its super block and bitmap blocks data .

Syscalls

To be able to have a simple functioning file system, we need to handle a minimum set of operations on files and directories. The functions are listed here in the order that I suggest you implement them in.

The syscalls need to return success or failure. Success is indicated by 0 and appropriate errors by the negation of the error code, as listed on the corresponding function's man page.

u_fs_getattr	Description:	This function should look up the input path to determine if it is a directory or a file. If it is a directory, return the appropriate permissions. If it is a file, return the appropriate permissions as well as the actual size. This size must be accurate since it is used to determine EOF and thus read may not be called.
	UNIX Equivalent:	man -s 2 stat
	Return values:	0 on success, with a correctly set structure -ENOENT if the file is not found
u_fs_readdir	Description:	This function should look up the input path, ensuring that it is a directory, and then list the contents. To list the contents, you need to use the filler() function. For example: filler(buf, ".", NULL, 0); adds the current directory to the listing generated by ls -a In general, you will only need to change the second parameter to be the name of the file or directory you want to add to the listing.
	UNIX Equivalent:	man -s 2 readdir However it's not exactly equivalent
	Return values:	0 on success -ENOENT if the directory is not valid or found
u_fs_mkdir	Description:	This function should add the new directory to the root level, and should update the .directories file appropriately.
	UNIX	man -s 2 mkdir

Equivalent:	
Return values:	0 on success - ENAMETOOLONG if the name is beyond 8 chars - EPERM if the directory is not under the root dir only - EEXIST if the directory already exists

u_fs_rmdir

Description:	Deletes an empty directory
UNIX Equivalent:	man -s 2 rmdir
Return values:	0 read on success - ENOTEMPTY if the directory is not empty - ENOENT if the directory is not found - ENOTDIR if the path is not a directory

u_fs_mknod

Description:	This function should add a new file to a subdirectory, and should update the .directories file appropriately with the modified directory entry structure.
UNIX Equivalent:	man -s 2 mknod
Return values:	0 on success - ENAMETOOLONG if the name is beyond 8.3 chars - EPERM if the file is trying to be created in the root dir - EEXIST if the file already exists

u_fs_write

Description:	This function should write the data in buf into the file denoted by path, starting at offset.
UNIX Equivalent:	man -s 2 write
Return values:	size on success - EFBIG if the offset is beyond the file size (but handle appends)

u_fs_read

Description:	This function should read the data in the file denoted by path into buf, starting at offset.
UNIX Equivalent:	man -s 2 read
Return values:	size read on success - EISDIR if the path is a directory

u_fs_unlink

Description:	Delete a file
UNIX Equivalent:	man -s 2 unlink
Return values:	0 read on success - EISDIR if the path is a directory - ENOENT if the file is not found

u_fs_open

This function should not be modified, as you get the full path every time any of the other functions are called.

u_fs_flush

This function should not be modified.

`u_fs_truncate` This function should not be modified.

Building and Testing

Your source files should be included as part of the `Makefile` in your directory, so building your changes is as simple as typing `make`.

One suggestion for testing is to launch a FUSE application with the `-d` option (`./u_fs -d testmount`). This will keep the program in the foreground, and it will print out every message that the application receives, and interpret the return values that you're getting back. Just open a second terminal window and try your testing procedures. Note if you do a **CTRL+C** in this window, you may not need to unmount the file system, but on crashes (transport errors) you definitely need to.

Your first steps will involve simply testing with `ls` and `mkdir`. When that works, try using `echo` and redirection to write to a file. `cat` will read from a file, and you will eventually even be able to launch `pico` on a file.

Remember that you may want to delete your `.directories` or `.disk` files if they become corrupted. You can use the commands `od -x` to see the contents in hex of either file, or the command `strings` to grab human readable text out of a binary file.

Notes and Hints

- The root directory is equivalent to your mount point. The FUSE application does not see the directory tree outside of this position. All paths are translated automatically for you.
- `sscanf(path, "%[^/]/%[^.].%s", directory, filename, extension);`
- Your application is part of userspace, and as such you are free to use whatever parts of C Standard Library you wish, including the file handling functions.
- Start early; this one will take a bit of time.
- Remember to always close your disk and directory files after you open them in a function. Since the program doesn't terminate until you unmount the file system, if you've opened a file for writing and not closed it, no other function can open that file simultaneously.
- Remember to open your files for binary access.

Requirements and Submission

You need to submit:

- Your well-commented source program
- A technic report within 20 pages

You can write your technic report either by English or Chinese. Bellow is a skeleton of technic report.

操作系统课程设计

题目：一个用户级文件系统的设计

姓名：

学号：

班级：

电话：

电子邮件：

Qq：

- 1、 课程设计的主要目的
- 2、 相关的技术背景
- 3、 主要思想和技术路线
- 4、 测试结果
- 5、 源代码的目录结构及存放位置
- 6、 运行环境